

RESPONSE TO SELECTION CRITERIA

Candidate: Troy Latter **Position:** Technology Manager – AI Platforms & Operations **Organisation:** UNSW Sydney, Division of Operations **Date:** 6 August 2026

1. Relevant tertiary qualification in Information Technology, Engineering, Computer Science, Data/AI, or equivalent

I hold a Master of Commerce (Marketing) from the University of New South Wales and a Bachelor of Economics (Law / Industrial Relations) from the University of Wollongong. While my formal qualifications are not in computer science, they are reinforced by more than 20 years of technical leadership across enterprise architecture, cloud, AI, data, service management and large-scale transformation. I am also a UNSW alumnus, which gives me a direct appreciation of the University's role as both a complex institution and a public trust.

My formal credentials include AWS Certified Solutions Architect – Professional, TOGAF Foundation, ITIL v3, IT4IT Foundation, PRINCE2 Practitioner and Prosci ADKAR Change Management, together with cloud and AI platform micro-credentials across AWS, Microsoft Azure, Google Cloud and Oracle. These credentials align directly with the capability profile reflected in my CV: enterprise architecture, engineering leadership, AI platform strategy, operational excellence, service management, governance and organisational change.

The strongest evidence of equivalence is the scope of the roles I have held. At Unisys, I led regional technology strategy and a solutions engineering capability of more than 15 people across six countries. At AWS, I advised more than 50 CIOs and technology leaders across Australia and ASEAN and led architecture reviews, technical discovery, proof-of-value assessment and implementation planning for large-scale cloud, data and AI environments. At Oracle and HPE, I worked across Home Affairs, Defence, Finance, health and critical infrastructure, where technical decisions had to satisfy stringent requirements for security, sovereignty, continuity and operational supportability.

This combination of academic qualifications, recognised technical credentials and long-term executive practice gives me the breadth required for the role. At UNSW, I would bring not only technical knowledge, but also the ability to translate institutional strategy into funded, sequenced and supportable technology capability across research, teaching and operations.

2. Demonstrated experience leading AI/ML or advanced technology delivery and operations teams in complex organisations, with accountability for outcomes, service performance and risk management at enterprise or multi-business-unit scale

I have repeatedly led advanced technology teams and cross-functional engineering groups in complex, regulated and geographically distributed environments. My experience covers regional engineering leadership, enterprise architecture, AI platform evaluation, cloud and data delivery, governance, service maturity, partner management and executive accountability. The strongest example is my role as APAC CTO and Head of Strategic Foresight at Unisys, where I led regional technology strategy and a solutions engineering capability of more than 15 people across six countries supporting government, financial services, telecommunications and critical-infrastructure clients.

At Unisys, I focused on creating repeatable delivery capability rather than isolated technical wins. I established engineering playbooks, solution patterns, architecture approaches and executive briefing frameworks that improved consistency and readiness across the region. I worked across product, architecture, sales and delivery to prioritise platform capability, customer outcomes and investment. I also directed secure LLM and agentic AI pilots for procurement, quoting, compliance and service workflows, ensuring that explainability, privacy, cybersecurity, orchestration controls and operational supportability were considered before any use case progressed.

My AWS experience strengthened this at a larger ecosystem scale. As Principal Solutions Architect and Global Capability Adviser, I led cross-functional technical teams through architecture reviews, discovery, solution validation and implementation planning for large-scale cloud, data and AI environments. I advised more than 50 CIOs and senior technology leaders, delivered more than 100 executive and architecture sessions, and supported initiatives including Cross River Rail, disaster recovery and COVID-response architectures. The key capability was not only technical design, but ensuring proposed solutions were secure, resilient, supportable and understood by the teams that would ultimately operate them.

My operational discipline was built earlier at HPE, where I applied ITIL, IT4IT, SIAM and ServiceNow practices across Defence, Finance and Human Services. I supported incident, problem, service transition, operational readiness and continuous-improvement disciplines in complex multi-vendor environments. This experience is reflected in the operational and governance skills set out in my CV: service management, SLAs/SLOs, observability, capacity planning, enterprise risk, Responsible AI and architecture governance.

At UNSW, I would bring this combination of engineering leadership and operational discipline to the internal AI technology delivery and platform opera-

tions teams. My focus would be on clear standards, strong technical leadership, transparent accountability, effective partner management, safe pilot-to-scale decisions and service performance that is measured and improved over time.

3. Proven ability to translate strategic goals into technical roadmaps and operational plans, including platform maturity, delivery standards and measurable service targets that support organisation-wide transformation and capability uplift

Translating strategy into executable roadmaps is one of the most consistent themes in my career. My approach is to connect strategic intent with architecture, operating model, investment sequencing, governance, delivery standards and organisational capability so that a roadmap survives contact with delivery reality. This capability is reflected throughout my CV in AI product and platform strategy, technology lifecycle management, proof-of-concept to scale, portfolio governance, capability uplift and change leadership.

At Unisys, I shaped AI, cloud and service-transformation roadmaps across Microsoft, AWS, Google and IBM ecosystems. I evaluated platform options against security, governance, operational supportability, customer value and strategic fit, then used those findings to guide investment priorities, product and platform direction and the progression of AI pilots. I worked across product, architecture, sales and delivery so that roadmaps were not purely technical artefacts; they became shared operating plans with clearer sequencing, ownership and decision points.

At Tech 4 Humanity, I developed product and platform roadmaps spanning AI orchestration, consent, workforce augmentation, trusted data use and organisational capability. These were structured around staged progression from pilot to proof of value to scale, with governance, integration, supportability and human oversight built in from the beginning. That work reinforced my belief that AI roadmaps must include the technical platform, the operating model, the people capability and the governance system together, rather than treating them as separate workstreams.

At AWS and Oracle, I applied the same method in major government and infrastructure environments. This included Cross River Rail, Home Affairs, immigration and identity modernisation, national health, smart infrastructure and resilience programs. In each case, I translated executive priorities into practical sequencing, architecture options, risk controls and capability-development plans. At HPE, I then connected these roadmaps to operational-readiness, service-management and business-as-usual transition requirements.

At UNSW, I would apply this experience to align the AI Ecosystem roadmap with platform maturity, delivery standards, service expectations, capacity, risk, governance and capability uplift. The result would be a roadmap that gives

the Head of AI Ecosystem, Product Manager, Enterprise Architect, engineering teams and institutional stakeholders a shared view of priorities, dependencies, measures and decisions.

4. Strong knowledge of AWS and Azure AI/ML and integration services, including enterprise deployment patterns

My experience spans architecture, platform evaluation, adoption strategy, governance and integration across Microsoft, AWS, Google, IBM and Oracle environments. I have deliberately positioned this experience accurately: strong in enterprise platform evaluation, architecture and adoption leadership, with direct depth in AWS and practical knowledge of Microsoft AI and integration services. This aligns with the technology portfolio in my CV, which includes Azure AI, Azure OpenAI, Microsoft Copilot, Power Platform, AWS AI/ML, SageMaker, AWS Bedrock familiarity, API Management, serverless architecture, data platforms and multi-cloud integration.

At Unisys, I evaluated AI platforms across Microsoft, AWS, Google and IBM ecosystems against enterprise criteria such as strategic fit, security, governance, supportability, integration readiness and customer value. This included work involving Azure AI, Azure OpenAI, Microsoft Copilot, Power Platform and multi-cloud integration patterns. My responsibility was to help determine which options were appropriate for particular use cases, what controls were required and how platform decisions would affect architecture, operations and long-term supportability.

My AWS experience is deeper and spans more than three years as Principal Solutions Architect and Global Capability Adviser. I led architecture reviews, technical discovery, proof-of-value assessment and implementation planning using AWS cloud, data, AI/ML, serverless and integration services. I shaped enterprise roadmaps involving SageMaker, API Gateway, Lambda, data platforms and resilient cloud patterns, and I supported customers in understanding the implications of scale, security, cost, reliability and operational readiness.

Across Oracle and HPE, I worked with hybrid, sovereign and bandwidth-constrained environments where integration and security architecture were critical. These roles required me to balance legacy-system constraints, identity, data movement, service continuity, sovereignty and operational support. That experience is directly relevant to a university environment, where AI platforms must integrate with existing systems without fragmenting the architecture or creating support burdens.

At UNSW, I would bring an objective, multi-cloud perspective. I can evaluate Azure and AWS options against architecture, security, governance, integration, cost and serviceability, and work with the Enterprise Architect and Product Manager to ensure the selected patterns are technically coherent, supportable and aligned with the institutional roadmap.

5. Proven ability to manage operational performance and service quality, including SLAs/SLOs, observability, incident/problem management, capacity planning and continuous service improvement

Operational excellence is a core part of my leadership approach. My strongest formal foundation was built at Hewlett Packard Enterprise, where I applied ITIL, IT4IT, SIAM and ServiceNow practices across Defence, Finance and Human Services. Since then, I have carried the same service-management discipline into architecture, cloud, AI and executive advisory roles. My CV reflects this directly through service management, SLAs/SLOs, observability, incident and problem management, capacity planning, operational monitoring and continuous improvement.

At HPE, I supported incident, problem, service transition, operational readiness and continual-improvement disciplines across complex multi-vendor environments. I co-developed business cases, operating models, architecture roadmaps and change strategies, and advised senior executives on service maturity, benefit realisation and technology risk. The central lesson was that even the strongest architecture fails if it cannot be supported, monitored, governed and transitioned into business-as-usual operations.

At AWS, I applied operational excellence, reliability, security and performance principles to large-scale cloud environments. I supported resilience and disaster-recovery architectures, including Cross River Rail and COVID-response contexts where continuity and rapid response were critical. I led architecture reviews that considered not only design intent, but also availability, failure modes, recovery, operational ownership and the capability of customer teams to run the environment after implementation.

At Unisys, I embedded operational supportability into AI platform evaluation and pilot decisions. This included asking whether ownership was clear, whether escalation and support paths existed, whether security and governance controls could be maintained, and whether the use case could move beyond demonstration into a repeatable service. At Tech 4 Humanity, I extended this thinking into trust-oriented design patterns covering logging, audit trails, human escalation, provenance, runtime controls and policy alignment.

At UNSW, I would apply this discipline through clear service definitions, tiered service expectations, incident and problem governance, operational reporting, capacity planning, vendor accountability and continual improvement. The objective would be to ensure AI services are not only innovative, but reliable, transparent, supportable and able to mature as institutional demand grows.

6. Demonstrated ability to evaluate AI products and solutions, balancing value, risk, compliance and operational supportability while informing enterprise investment decisions

Evaluating emerging technology is one of the clearest themes running through my career. I have repeatedly been asked to help senior leaders decide not simply whether a platform or solution is innovative, but whether it is strategically relevant, technically viable, governable, supportable and worthy of enterprise investment. That work has required commercial judgement, architecture depth, risk awareness and the confidence to recommend that a pilot be stopped or reworked when the evidence did not support scale.

At Unisys, as APAC CTO and Head of Strategic Foresight, I evaluated AI platforms and services across Microsoft, AWS, Google and IBM ecosystems against enterprise criteria including security, privacy, strategic fit, integration complexity, operational supportability, vendor viability and customer value. I directed secure LLM and agentic AI pilots for procurement, quoting, compliance and service workflows, and worked across product, architecture, sales and delivery to determine which use cases should progress, be redesigned or be stopped. The same evidence base informed internal investment priorities, customer roadmaps and a more coherent regional AI capability rather than a collection of disconnected experiments.

At AWS, I partnered with more than 50 CIOs and technology leaders to assess proof-of-value initiatives in cloud, data and AI/ML. My role was to move the conversation beyond a technical demonstration and test whether the initiative could satisfy enterprise requirements for security, resilience, cost, operating ownership and long-term support. I applied the same discipline at Oracle in major public-sector programs spanning immigration, identity, health and smart infrastructure, where technology options had to be assessed against sovereignty, continuity, legislative obligations and the consequences of failure. At Tech 4 Humanity, I have extended that approach into practical pilot-to-scale roadmaps for organisations in education, aged care, transport and the public sector.

My evaluation model is consistent: clarify the intended outcome; test architecture and integration fit; assess security, privacy and Responsible AI obligations; examine service ownership, operating model and whole-of-life cost; then make a scale, remediate or stop decision. This draws directly on the skills reflected in my CV: product and platform strategy, architecture governance, enterprise risk, service management, investment planning and executive advisory. It also reflects the evidence in my career examples, where the strongest outcomes came from improving decision quality rather than encouraging technology adoption for its own sake.

At UNSW, I would apply this discipline through a clear pilot-to-scale framework aligned with the AI Ecosystem roadmap. Product value, architectural coherence,

governance compliance and operational readiness would be assessed together, not sequentially after investment has already been committed. This would help reduce pilot fatigue, avoid avoidable technical debt and ensure that enterprise funding is directed toward AI services that can deliver measurable institutional value while remaining secure, explainable and supportable.

7. Excellent stakeholder management and collaboration skills across technical and non-technical teams, including communicating complex technical topics clearly to senior audiences and influencing executive decision-making

My career has been built around working at the point where technical complexity, institutional priorities and executive accountability meet. I am comfortable moving between engineering teams, architects, product leaders, cybersecurity and privacy specialists, commercial stakeholders, operational teams, executives and public-sector decision-makers. My contribution is not simply to explain technology clearly; it is to help different groups reach an informed decision despite competing priorities, different vocabularies and uneven levels of technical confidence.

At AWS, I advised more than 50 CIOs and senior technology leaders across Australia and ASEAN and delivered more than 100 executive briefings, architecture sessions and technical workshops. Those engagements covered cloud modernisation, AI/ML adoption, resilience, data platforms and enterprise architecture. The practical value came from framing choices in terms of business outcomes, risk, investment, sequencing and capability rather than presenting a list of services. I regularly worked with customer executives, internal specialists, partners and delivery teams to turn strategic ambition into a common roadmap and to build the confidence required for a decision.

At Oracle, I worked with government and national CIO stakeholders across Home Affairs, immigration, identity, healthcare and smart-infrastructure programs. These environments required alignment across policy, procurement, legal, security, architecture and implementation communities, often across multiple jurisdictions. At Unisys, I connected regional engineering, product, architecture, sales and delivery stakeholders around AI and cloud roadmaps and customer priorities. Earlier, in Australian Government transformation work, I led change and stakeholder engagement across immigration reform spanning 65 countries, where cultural context, local operating conditions and distributed authority made one-size-fits-all communication ineffective.

The skills supporting this capability are visible throughout my CV: executive advisory, architecture leadership, product partnership, organisational change, stakeholder leadership and capability uplift. My approach is to listen first, understand what each stakeholder is accountable for, then translate the issue into shared decision language. I do not hide uncertainty or oversell AI. I present trade-offs, identify what must be true for a proposal to work and make the deci-

sion path explicit. This builds credibility with technical teams because the detail is respected, and with executives because the discussion remains connected to institutional risk and value.

At UNSW, this capability would be central to the role. The Technology Manager will need to work across the Head of AI Ecosystem, Product Manager, Enterprise Architect, engineering and operations teams, cybersecurity, privacy, faculties, researchers, professional services, vendors and senior university leadership. I would bring the ability to build alignment without relying solely on formal authority, translate GenAIOps and LLMOps issues into institutional consequences, and maintain trust when priorities or technical realities require difficult decisions.

8. Excellent communication and reporting skills, including preparing operational summaries and updates for senior stakeholders

Clear communication and disciplined reporting have been essential in every senior role I have held. I believe executive reporting should make the state of delivery easier to understand, not bury it under activity measures or technical jargon. A useful report should show what changed, what is performing, what is at risk, what decisions are required and what action will follow. That principle has guided my executive briefings, architecture reviews, business cases, operating-model work and strategic roadmaps.

At AWS, I delivered more than 100 executive briefings, architecture sessions and technical workshops to CIOs, technology leaders, partners and delivery teams. I prepared architecture options, risk assessments, roadmap recommendations and investment narratives that enabled senior stakeholders to make informed decisions about cloud, AI/ML, resilience and modernisation. At Unisys, I created reusable executive briefing frameworks and technical playbooks that improved the consistency and quality of regional decision support. I also advised CIOs and executive teams on AI opportunity, architecture choices, investment trade-offs and long-term technology direction.

At HPE, I co-developed business cases, investment models, architecture roadmaps, operational-readiness assessments and change strategies for major public-sector programs across Defence, Finance and Human Services. At Oracle, I communicated complex architecture and sequencing decisions across government, health, immigration and smart-infrastructure environments where senior stakeholders needed clarity on security, sovereignty, continuity and long-term supportability. My broader program and change background also means I understand the importance of reporting progress, risks, dependencies, benefits and stakeholder impacts as one integrated management picture.

This capability aligns directly with the communication, executive advisory, service management, portfolio governance and stakeholder leadership skills pre-

sented in my CV. My style is concise, evidence-based and candid. I distinguish between fact, assumption and recommendation, and I avoid creating false certainty around emerging technologies. For operational reporting, I would expect a consistent rhythm covering service performance, incidents and problems, capacity, risk, cost, roadmap delivery, vendor performance and capability maturity, with detail scaled to the audience.

At UNSW, I would establish an executive reporting rhythm that supports the Head of AI Ecosystem and wider Division of Operations leadership without creating unnecessary administrative load. Weekly operational summaries would highlight immediate performance and risk; monthly platform-health reviews would address service, cost, incidents, capacity and roadmap progress; and quarterly reviews would connect operational evidence to strategic investment and capability decisions. The objective would be transparent decision support and institutional confidence, not reporting for its own sake.

9. Knowledge of health and safety responsibilities and commitment to attending relevant health and safety training

I understand that health and safety is a leadership responsibility and a prerequisite for sustainable delivery. In technology and transformation environments, this includes physical safety, but it also includes psychosocial hazards such as excessive workload, poor role clarity, persistent operational pressure, unmanaged change, moral concern about technology outcomes and cultures where staff do not feel safe to raise risks. Senior leaders are responsible for recognising these conditions early and ensuring they are managed through normal planning and governance.

Across large public-sector transformation programs, including immigration reform spanning 65 countries, I worked in environments involving sustained delivery pressure, operational disruption, diverse workforces and significant public impact. My response has been to use realistic sequencing, workload management, role clarity, structured escalation, change planning and active stakeholder engagement to reduce unnecessary stress and avoid transferring delivery pressure onto teams without support. My service-management background at HPE also reinforced the value of clear incident processes, defined ownership and disciplined transitions rather than relying on individual heroics.

For AI and platform operations, I recognise additional risks arising from on-call expectations, rapidly evolving tools, uncertain accountability and the consequences of technology decisions affecting staff, students or the public. My leadership approach is to include people risk within operational governance, not treat it as separate from delivery. This aligns with the risk, service management, organisational change, incident/problem management and leadership capabilities reflected in my CV. I am committed to complying with UNSW's health and safety framework, attending all required training and ensuring hazards are identified, assessed, controlled and reviewed across the work for which I am

accountable.

Closing Statement

The Technology Manager – AI Platforms & Operations role requires a leader who can connect engineering leadership, platform strategy, service operations, governance, investment judgement and organisational capability. My career has consistently operated at that intersection across Unisys, AWS, Oracle, HPE, Tech 4 Humanity and complex public-sector environments.

I would bring UNSW a combination of regional engineering leadership, AI platform evaluation, enterprise architecture, Responsible AI governance, service-management maturity, executive communication and large-scale change experience. My focus would be practical and institutional: build stronger teams, establish clear standards, improve operational visibility, make better pilot-to-scale decisions and help UNSW convert AI ambition into secure, scalable and supportable capability.